

CSC420 Assignment 2

Owner's Information

First Name: Kwan Kiu

Last Name: Choy

Student No: 1005005879

UtorId: choykwan

Theoretical Part

[Question 1](#)

[Question 2](#)

[Question 3](#)

[Question 4](#)

[Question 6](#)

Implementation Part

[Task 1: Data Preprocessing](#)

[Task 2: Testing learning rates and Optimizer](#)


[Task 3: Testing the Number of Hidden Units](#)

[Task 4: Testing the Number of Layers](#)

[Task 5: Testing the Dropout](#)

Theoretical Part

Question 1



Let z_1, z_2 be the output of the 1st and 2nd layer.
 Let W_1, W_2 be the weight matrices of size $m \times n$ matrix
 where n is the number of the inputs and m is the
 number of the units of the hidden layers. Let b_1, b_2
 be the biases of each hidden layer. Let $f(x) = ax + C$ be the
 linear activation function and I_1, I_2 be the output of
 the activation function

$$z_1 = W_1 X + b_1$$

$$I_1 = f(z_1) = az_1 + C = aW_1 X + ab_1 + C = aW_1 X + b_1'$$

$$z_2 = aW_2 W_1 X + W_2 b_1' + b_2$$

$$\begin{aligned} I_2 &= a(aW_2 W_1 X + W_2 b_1' + b_2) + C \\ &= a^2 W_2 W_1 X + aW_2 b_1' + ab_2 + C \\ &= \underline{a^2 W_2 W_1} X + \underline{b_2'} \end{aligned}$$

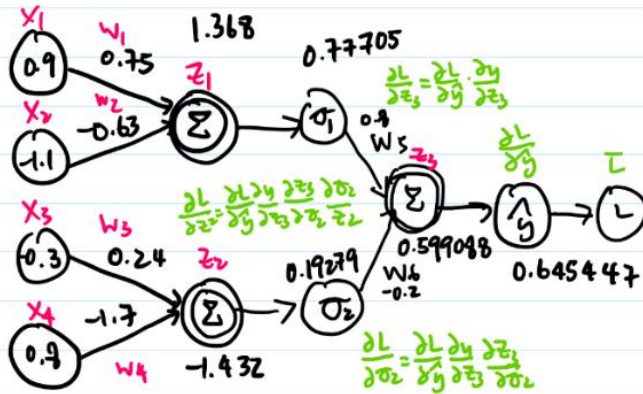
As a result, we have shown that I_2 can be
 expressed as $W'X + b_2'$ which means that it can
 be expressed as a linear function of X . The slope
 of the linear function is the W' matrix.

This aligns with our understanding since we are
 only applying a linear function on the weighted
 sum between the weights and the inputs at each
 hidden layer. We know that the composition of 2
 linear function is just another linear function. No
 matter how many times we apply $f(x)$,
 the result will still be a linear function of inputs.

No matter what k is, z_1, z_2 can always be expressed
 as a linear function of the inputs, X . This means that
 back propagation will not have any effect since
 the derivative is constant.

Question 2

2)



Equations:

- ① $L = ||0.5 - \hat{y}||^2$
- ② $\hat{y} = \frac{1}{1 + e^{-z_3}}$
- ③ $z_3 = w_5 \sigma_1 + w_6 \sigma_2$
- ④ $\sigma_1 = \frac{1}{1 + e^{-z_1}}$
- ⑤ $\sigma_2 = \frac{1}{1 + e^{-z_2}}$
- ⑥ $z_2 = w_3 X_3 + w_4 X_4$

Partial derivatives

$$\bar{L} = 1$$

$$\frac{\partial L}{\partial \hat{y}} = -2 ||0.5 - \hat{y}||$$

$$\frac{\partial \hat{y}}{\partial z_3} = \frac{e^{-z_3}}{(1 + e^{-z_3})^2}$$

$$\frac{\partial z_3}{\partial \sigma_1} = w_5$$


$$\frac{\partial \sigma_2}{\partial z_2} = \frac{e^{-z_2}}{(1 + e^{-z_2})^2}$$

$$\frac{\partial z_2}{\partial w_3} = X_3$$


$$\begin{aligned} \frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial \bar{L}} \cdot \frac{\partial \bar{L}}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_3} \cdot \frac{\partial z_3}{\partial \sigma_2} \cdot \frac{\partial \sigma_2}{\partial z_2} \cdot \frac{\partial z_2}{\partial w_3} \\ &= -2 ||0.5 - \hat{y}|| \frac{e^{-z_3}}{(1 + e^{-z_3})^2} w_6 X_3 \frac{e^{-z_2}}{(1 + e^{-z_2})^2} \\ &= -2 ||0.5 - 0.64545|| \frac{e^{-0.599098}}{(1 + e^{-0.599098})^2} (-0.2) (-0.3) \frac{e^{-1.432}}{(1 + e^{-1.432})^2} \\ &= 6.2158 \times 10^{-4} // \end{aligned}$$

Question 3

3)



Output of C has 20 channels



1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36



1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Output of convolution

16 max pooling operations

In max pooling, there are 9 pixels per filter so each pooling accords for 9 FLOPS

After convolution, we have an image of 6x6. We do convolution 36 times, and there are 20 filters in total.

★ Each convolution:

→ 50x16 multiplications

→ 50x16-1 addition

Total number of operations in convolution layer:

$$20 \times 36 \times (50 \times 16 + 50 \times 16 - 1) = 1151280.$$

Total number of operations in max pooling layer.

$$16 \times 8 \times 20 = 2560$$

★ Total number of operations in the network without bias:

$$1151280 + 2560 = 1153840 //$$

★ Total number of operations in network with bias:

Each convolution operation

- 50x16 multiplications

- 50x16 additions

$$20 \times 36 \times (50 \times 16 + 50 \times 16) + 16 \times 8 \times 20 = 1154560 //$$

Question 4

4) First layer: Input $\rightarrow C_1$

We know that the filter size can be computed by the size of the input tensor W_1 and the output tensor W_2 .

$$W_2 = (W_1 - F) / S + 1 \quad \text{F-filter size } S\text{-stride size}$$

$$28 = (32 - F) / 1 + 1$$

$$27 = 32 - F$$

$$F = 5$$

The number of trainable parameters per feature map in C_1 :

total number pixels in filter + bias

$$= 5 \times 5 + 1 = 26$$

We have 6 feature maps in C_1 , so total for this layer:

$$26 \times 6 = 156$$

Second layer: $C_1 \rightarrow S_2$

Assuming that the subsampling operation is performed by a max pooling or an average pooling layer, there would not be any trainable parameters in this layer.

Third layer: $S_2 \rightarrow C_3$

using formula before,

$$W_2 = (W_1 - F) / S + 1$$

$$10 = (14 - F) / 1 + 1$$

$$9 = 14 - F$$

So the filter needed $F = 5$

in the convolution operation is a 5×5 filter.

Each of the feature map would have its own bias, and therefore the total number of trainable parameters would be

$$6 \times (5 \times 5 + 1) \times 16$$

$$= 6 \times 26 \times 16 = 2496$$

Fourth layer: $C_3 \rightarrow S_4$

Assuming that the subsampling operation is performed by a max pooling or an average pool (fixed pooling method) there are no trainable parameters.

Fifth layer: $S_4 \rightarrow C_5$

We know the filter is 5×5 and thus there's going to be only 1 bias term for each feature map in the C_5 layer

$$(16 \times 5 \times 5 + 1) \times 20 = 48120$$

Sixth layer: $C_5 \rightarrow F_6$

There's 1 bias for each pixel in F_6 layer, so the total number of trainable parameters:

$$84 (120 + 1) = 10164$$

Seventh layer: $F_6 \rightarrow \text{Gaussian}$

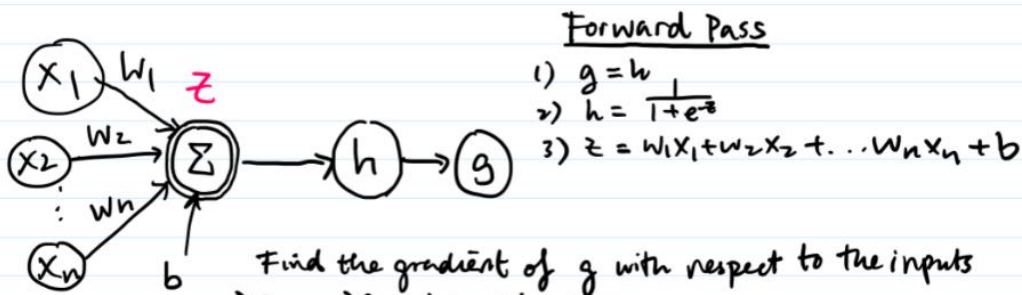
There's 1 bias for each pixel in the gaussian layer, so the number of trainable parameters:

$$10 (84 + 1) = 850$$

Total: $156 + 2496 + 48120 + 10164 + 850 = 61786$

Question 5

5)



Find the gradient of g with respect to the inputs

$$\begin{aligned} \frac{\partial g}{\partial x_1} &= \frac{\partial g}{\partial g} \cdot \frac{\partial g}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial x_1} \\ &= 1 \cdot 1 \cdot \frac{e^{-z}}{(1+e^{-z})^2} \cdot w_1 \\ &= \frac{w_1 e^{-z}}{(1+e^{-z})^2} \end{aligned}$$

Since we know that z can be found by using g using $z = -\ln(\frac{1}{g}-1)$.

$$\rightarrow \text{So } \frac{\partial g}{\partial x_1} = \frac{w_1 e^{\ln(\frac{1}{g}-1)}}{(1+e^{\ln(\frac{1}{g}-1)})^2}$$

$$\begin{aligned} \text{[2]} \quad \frac{\partial g}{\partial x_2} &= \frac{\partial g}{\partial g} \cdot \frac{\partial g}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial x_2} \\ &= 1 \cdot 1 \cdot \frac{e^{-z}}{(1+e^{-z})^2} w_2 \\ &= \frac{w_2 e^{-z}}{(1+e^{-z})^2} \\ &= \frac{w_2 e^{\ln(\frac{1}{g}-1)}}{(1+e^{\ln(\frac{1}{g}-1)})^2} \end{aligned}$$

So $\forall i \in \mathbb{N}, 1 \leq i < n$


$$\begin{aligned} \frac{\partial g}{\partial x_i} &= \frac{\partial g}{\partial g} \cdot \frac{\partial g}{\partial h} \cdot \frac{\partial h}{\partial z} \cdot \frac{\partial z}{\partial x_i} \\ &= \frac{w_i e^{\ln(\frac{1}{g}-1)}}{(1+e^{\ln(\frac{1}{g}-1)})^2} \end{aligned}$$

As a result, we have shown that the gradient of g does not depend on any of the input x_i .

Question 6

6a) Logistic Regression: $\frac{1}{1+e^{-wx+b}}$

Tanh(x): $\frac{1-e^{-2x}}{1+e^{-2x}}$



The shape of the two functions are very similar but the range of Tanh(x) function is from -1 to 1 whereas it is 0 to 1 for the logistic.

6b) $\tanh(x) = \frac{1-e^{-2x}}{1+e^{-2x}}$

Let $f(x)$ represents the logistic regression function such that $f(x) = \frac{1-e^{-x}}{1+e^{-x}}$

$$\begin{aligned}\tanh(x) &= (1-e^{-2x}) \cdot \frac{1}{(1+e^{-2x})} \\ &= (1-e^{-2x}) \cdot f(2x) \\ &= (2 - \frac{1}{f(2x)}) \cdot f(2x) \\ &= 2f(2x) - 1\end{aligned}$$

$$\begin{aligned}\frac{d}{dx} \tanh(x) &= \frac{d}{dx} (2f(2x) - 1) \\ &= 2f'(2x) \cdot 2 \\ &= 4f'(2x)\end{aligned}$$

Derivative of logistic regression:

$$\begin{aligned}f(x) &= \frac{1}{1+e^{-x}} \\ f'(x) &= \frac{e^{-x}}{(1+e^{-x})^2} \\ &= [f(x)-1]f'(x) \\ &= f(x) - f^2(x)\end{aligned}$$

Continuing with $\frac{d}{dx} \tanh(x)$,

$$\begin{aligned}&= 4f'(2x) \\ &= 4(f(2x) - f^2(2x)) \\ &= 4f(2x) - 4f^2(2x)\end{aligned}$$

As a result, we know that the gradient of the tanh(x) can be expressed as a function of the logistic regression functions.

6c) Tanh(x):

↳ It is symmetric around 0. If the mean of the inputs is 0, then it gives network to adjust the weights more flexibly as they don't have to all increase or all decrease together, which makes the backpropagation process more efficient. It centers data so that it's easier to train.

Logistic(x)

The range of the tanh(x) is larger than that of the logistic regression, as shown in (a). As a result, leads to larger derivatives (due to the shape of the curve and the large range), this means that convergence will be reached faster.

Implementation Part

Task 1: Data Preprocessing

```
def step_1_load_data(gaussian=False, normalized=True):
    """
    Preprocessed data and load the data such that trainloaders,
    validloaders
    and testloaders can be created for the rest of the codes to use.
    :param gaussian: Boolean to indicates whether we want gaussian blur or
    not
    in the data preprocessing stage
    :return:
    1) trainloader: loader to load training data
    2) validloader: Loader to load validation data
    3) testloader: loader to load testing data
    """
    # Define a transform to normalize the data
    transform = transforms.Compose([
        transforms.Grayscale(num_output_channels=1),
        transforms.ToTensor(),
        transforms.Normalize((0.5,), (0.5,))])
    if gaussian:
        transform = transforms.Compose([
            transforms.Grayscale(num_output_channels=1),
            transforms.GaussianBlur(3, sigma=0.7),
            transforms.ToTensor(),
            transforms.Normalize((0.5,), (0.5,))])
    if not normalized:
        transform = transforms.Compose([
            transforms.Grayscale(num_output_channels=1),
            transforms.ToTensor()])
    trainset = datasets.ImageFolder(root=TRAIN_DATA_PATH,
    transform=transform)
    trainloader = torch.utils.data.DataLoader(trainset, batch_size=64,
        shuffle=True)

    testset = datasets.ImageFolder(root=TEST_DATA_PATH,
    transform=transform)
    testloader = torch.utils.data.DataLoader(testset, batch_size=64,
        shuffle=True)
```



```

    validset = datasets.ImageFolder(root=VALID_DATA_PATH,
transform=transform)
    validloader = torch.utils.data.DataLoader(validset, batch_size=64,
                                                shuffle=True)

    return trainloader, testloader, validloader

```

My preprocessing pipeline:

- 1) Convert the images with 3 channels (RGB) into a single channel image (grayscale)
 - Using images with 3 channels would be problematic during training process, therefore I decided to convert it to grayscale images so that it is easier to train the network
- 2) Convert the images to tensors
 - This is a step that we must include in the preprocessing stage as stated in the assignment
- 3) Normalize the images with mean and std to be 0.5, which will normalize image in the range of [-1, 1]

	Training loss	Valid loss	Train accuracy	Valid accuracy
Without Normalization	1.53001837151	1.54884378623	0.934733333333	0.91499999904
With Normalization	1.5088002970	1.541611072540	0.9542	0.92399999952

Since both the training accuracies and validation accuracies improve with normalization, I decided to include that into my preprocessing pipeline.

Apart from the three steps mentioned in my data preprocessing state, I have also tried whether applying a gaussian blur to smoothen the image will improve the accuracy. The results are displayed below.

	Training loss	Valid loss	Train accuracy	Valid accuracy
With Gaussian	1.52371720536	1.54690929985	0.941666666666	0.918
Without Gaussian	1.50880029703	1.541611072540	0.9542	0.92399999952

Since, the validation loss is actually larger when a gaussian filter is applied, therefore I decided not to include gaussian blurring as part of the preprocessing stage.

Task 2: Testing learning rates and Optimizer

For this task, I have used 4 different optimizers and 5 different learning rates to test the training, validation and testing accuracies.

The 4 different Optimizers:

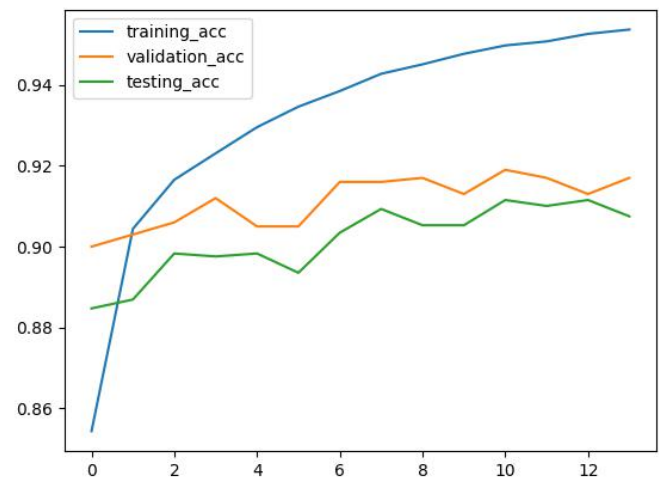

- 1) SGD
- 2) Adam
- 3) RMSProp
- 4) Adadelta

The 5 different learning rates:

- 1) 0.001
- 2) 0.01
- 3) 0.05
- 4) 0.1
- 5) 0.5


The table below shows the best results (lowest validation loss) for **each optimizer**.

- 1) Optimizer: Adadelta, Learning rate: 0.5, n_epochs = 14 (early stopped)



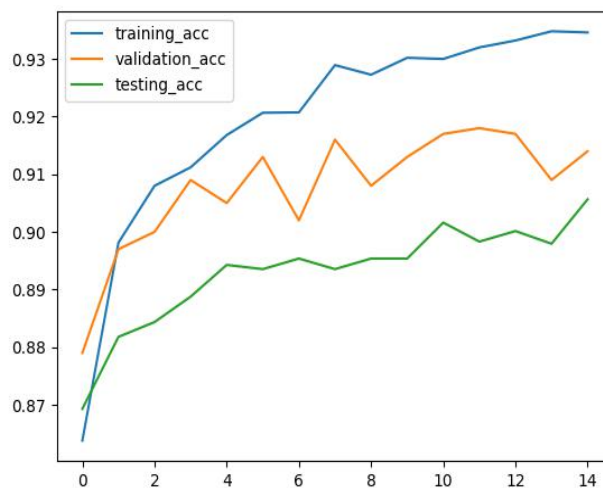
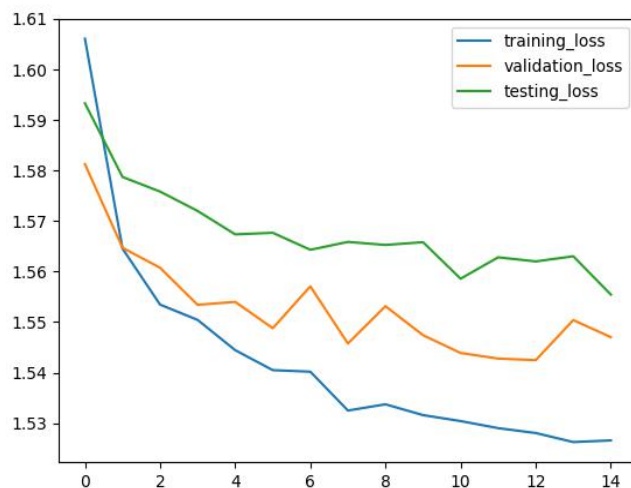
	Training	Validation	Test
Final Losses	1.509453068033854	1.546137075424194	1.555530561224479
Final Accuracies	0.9536666666666666	0.917000000953674	0.907488986434040

2) Optimizer: RMSProp, Learning rate: 0.5, n_epochs = 14 (early stopped)



	Training	Validation	Test
Final Losses	1.508800297037760	1.5416110725402832	1.552166826399412
Final Accuracies	0.9542	0.923999999523162	0.9115271661950278



3) Optimizer: Adam, Learning rate: 0.001, n_epochs = 15





	Training	Validation	Test
Final Losses	1.526564672915140	1.547003357887268	1.55542909678908
Final Accuracies	0.9346	0.914000000476837	0.90565345089516

4) Optimizer: SGD, Learning rate: 0.5, n_epochs= 14 (early stopped)



100



500



1000



Therefore, I ran the model with 100 units on the testing set, and the below are the results that I obtained.

Testing loss

1.5678207690495227

Testing Accuracies

0.8939060206455282


Task 4: Testing the Number of Layers

Using the best set of parameters in task 2 (optimizer: SGD, learning rate: 0.5), I ran experiments on both the one layer model (in task 2) and the new two layer model. The below are the summarized results.

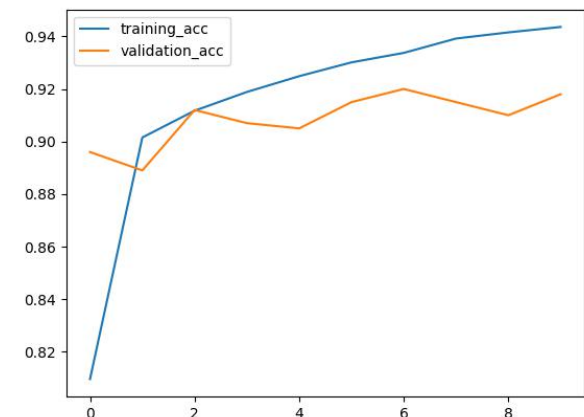


	Training Loss	Valid Loss	Training Accuracy	Valid Accuracy
Two Layer	1.518992037455	1.5446843318	0.94359999996821	0.9180000004768
One Layer	1.509279940032	1.5429654560	0.95326666666666	0.919

The two layer model and the one layer model have very similar validation loss. However, the one layer model has a slightly lower validation loss compared to the two layer models.

One Layer



Two Layer



	Testing loss	Testing Accuracies
Two Layer	1.5676022591219774	0.8957415567970836
One Layer	1.56034973119324	0.9034508075483045



From this table, we can see that the one layer model does in fact have a higher testing accuracy than the two layer model, which aligns with the training and validation results generated before.

Task 5: Testing the Dropout



Using the best set of parameters in task 2 (optimizer: SGD, learning rate: 0.5), I ran an experiment on the model we used in task 2 with a dropout layer. The below are the summarized results.

	Training loss	Valid loss	Training Accuracy	Valid Accuracy
With Dropout	1.509791118939	1.5525431972	0.95286666666348	0.911160058474
Without Dropout (From task 2)	1.508800297037	1.5416110725	0.9542	0.9239999999523

Without Dropout



With Dropout



Even though the final validation accuracy from the model with dropout is slightly higher than that of the model without dropout, the results from the plots are interesting. From the four plots above, it is very clear that the validation curve in the model with a dropout layer is smoother. There are less fluctuations (bumps) in the validation curve, and hence prevent the model from overfitting.